



CONFIDENTIAL CLIENT OPPORTUNITY • INDEPENDENT CANDIDATE MATERIALS

Learn before standardizing.

A discovery-led plan for a confidential global enterprise.

Entry principle: Do not begin with a technology roadmap. Begin with executive intent, the current portfolio, one representative workflow, decision rights, architecture reality, and the evidence needed to make the next investment decisions.

0–30

Read the system

Listen for

- Executive priorities and business-unit demand
- Current AI portfolio, pilots, platforms, and shadow activity
- Data, integration, identity, cloud, security, privacy, legal
- Team topology, vendor commitments, capacity, and decision rights

Produce

- Portfolio reality map
- Stakeholder and authority topology
- Architecture and governance baseline
- One representative workflow baseline
- Immediate risk / stop-work list only where necessary

31–60

Make decisions visible

Establish

- Enterprise AI intake and prioritization standard
- Fund / condition / narrow / hold / stop vocabulary
- Risk tiering and early review paths
- Build-vs-buy and exception decision records

Pilot the operating model

- Review two representative initiatives
- Map value, readiness, authority, architecture, adoption
- Set outcome and leading-indicator scorecard
- Confirm production ownership and review cadence

DAY-60 EXECUTIVE READOUT

Question	Evidence	Decision
Where is value being lost?	Portfolio and workflow baselines	Focus areas and WIP limits
Which controls arrive too late?	Review-path and rework analysis	Earlier authority and evidence gates
What should be standardized?	Repeated patterns and exception burden	Reference capabilities and owners
What should stop?	Weak ownership, economics, readiness, or fit	Hold / stop decisions with rationale

120-DAY ENTRY PLAN • DELIVER, EMBED & INSTITUTIONALIZE • PAGE CONTINUATION



CONFIDENTIAL CLIENT OPPORTUNITY • INDEPENDENT CANDIDATE MATERIALS

61-90

Deliver one scalable increment

Solution delivery

- Confirm reference architecture and exception rationale
- Define evaluation, acceptance, telemetry, cost, rollback
- Make human authority and escalation explicit
- Align delivery roles across engineers, architects, administrators, business owners, and vendors

Workflow embedding

- Redesign roles and standard work
- Prepare training and champion support
- Instrument retained use, friction, overrides, and task success
- Run a controlled release and rapid learning cadence

OPERATING CADENCE

Cadence	Purpose	Output
Weekly delivery and risk review	Resolve blockers, evidence gaps, ownership, and decisions	Updated conditions and escalations
Biweekly portfolio review	Sequence demand against capacity and readiness	Fund / condition / hold / stop decisions
Monthly architecture and governance learning	Convert repeated patterns and exceptions into standards	Reference pattern or revised guardrail
Monthly value and adoption review	Compare expected vs retained value in real work	Scale, revise, or retire recommendation

WHAT I WOULD NOT DO

No platform-first roadmap

Do not assume the tool portfolio defines the business portfolio.

No governance theater

Do not add review steps that cannot change a decision.

No pilot accumulation

Do not leave experiments alive because stopping feels politically difficult.

No adoption appendix

Do not treat changed work as a communications task after delivery.

Proof objective by day 90: one initiative has moved through an explicit yield system with accountable ownership, architecture and controls, accepted delivery, workflow adoption, instrumentation, and a decision on what to scale or change.

91-120

Institutionalize enterprise AI yield

Publish

- Enterprise AI strategy and sequenced roadmap
- Portfolio intake and prioritization operating standard
- Reference architecture and exception ownership
- Responsible-AI and human-authority model
- Executive scorecard and reporting cadence

Confirm

- Team accountabilities and capability gaps
- Vendor roles, economics, and knowledge transfer
- Business-unit ownership and champion network
- Operating support and lifecycle management
- Scale, hold, and retirement decisions



CONFIDENTIAL CLIENT OPPORTUNITY · INDEPENDENT CANDIDATE MATERIALS

DAY-120 DELIVERABLES

Deliverable	What it changes	Success signal
Enterprise AI roadmap	Connects investment to business priorities and capacity	Clear sequencing and explicit non-priorities
Portfolio operating model	Creates consistent intake, disposition, escalation, and learning	Faster decision time; fewer orphan pilots
Architecture and governance model	Moves controls early enough to influence design	Reduced late rework and clearer exceptions
Value and adoption scorecard	Balances outcome, flow, trust, use, and learning	Executive decisions change based on evidence
Capability and vendor plan	Clarifies internal ownership and external leverage	Knowledge transfer, transparent economics, durable support

DISCOVERY RISKS TO VALIDATE

Demand risk Executive urgency may bypass qualification and capacity.	Architecture risk Multiple providers and local solutions may create hidden integration and identity debt.
Ownership risk Business sponsorship may not translate into workflow and lifecycle accountability.	Adoption risk Success may be counted at go-live rather than retained behavior and outcome.

DAY-120 EXECUTIVE OPERATING STATE

Executive question	Visible evidence	Decision made easier
Which initiatives earn scarce capacity?	Value, accountable sponsor, readiness, architecture, adoption burden, and evidence plan	Fund, condition, narrow, hold, or stop
Where should controls become reusable?	Repeated risk, late rework, common integrations, and recurring exception patterns	Standardize a reference capability or retain a bounded exception
What should remain domain-specific?	Local context, consequence, regulation, workflow authority, and user reality	Preserve owned variation without fragmenting the enterprise
Has the solution changed work?	Roles, human authority, standard work, training, friction, overrides, and retained use	Embed, revise, or retire
Is value being retained?	Outcome, productivity, quality, risk, economics, trust, and learning speed	Scale, redesign, or stop investing

End state: executive leaders can see where attention and capital are going, why an initiative is advancing or stopping, who owns the work and risk, what evidence exists, and how portfolio learning is compounding.